

Data Mining

Prof. Dr. Stefan Kramer
Johannes Gutenberg-Universität Mainz

Acknowledgements

- E. Frank
- I. Witten
- ... and others

Outline

- k-Means
- Probability-based/model-based clustering: the EM algorithm

K-Means

Partitioning Algorithms: Basic Concept

- *Partitioning method*: construct a partition of a database D of n objects into a set of k clusters such that the sum of the squared distances is minimized

$$\sum_{m=1}^k \sum_{t_{mi} \in K_m} (C_m - t_{mi})^2$$

- Given a k , find a partition of k *clusters* that optimizes the chosen partitioning criterion
 - globally optimal: exhaustively enumerate all partitions
 - heuristic methods: *k-means* and *k-medoids* algorithms
 - *k-means* (MacQueen 67): each cluster is represented by the center of the cluster
 - *k-medoids* or PAM (partition around medoids) (Kaufman & Rousseeuw 87): each cluster is represented by one of the objects in the cluster

Goal of Clustering 1

- Total sum of pairwise distances:

$$T(C) = W(C) + B(C)$$

- Within cluster distances:

$$W(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \sum_{C(i')=k} d(x_i, x_{i'})$$

- Between cluster distances:

$$B(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \sum_{C(i') \neq k} d(x_i, x_{i'})$$

Goal of Clustering 2

- *Assuming K is fixed in this simple setting, maximizing $B(C)$ and minimizing $W(C)$ is the same!*
 - we assume real-valued data and use the *squared Euclidean distance*

$$d(x_i, x_{i'}) = \sum_{j=1}^p (x_{ij} - x_{i'j})^2 = \|x_i - x_{i'}\|^2$$

$$T = \frac{1}{2} \sum_{i=1}^N \sum_{i'=1}^N d(x_i, x_{i'}) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \left(\sum_{C(i')=k} d(x_i, x_{i'}) + \sum_{C(i') \neq k} d(x_i, x_{i'}) \right)$$

Goal of Clustering 3

- Thus, the goal is to minimize:

$$W(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \sum_{C(i')=k} \|x_i - x_{i'}\|^2 = \sum_{k=1}^K \sum_{C(i)=k} \|x_i - \bar{x}_k\|^2$$

$$\min_C \sum_{k=1}^K \sum_{C(i)=k} \|x_i - \bar{x}_k\|^2$$

$$\min_{C, \{m_k\}_1^K} \sum_{k=1}^K \sum_{C(i)=k} \|x_i - m_k\|^2$$

- *Alternating optimization procedure...*

K-Means Algorithm


Given data points $D = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$, find K clusters $\{C_1, \dots, C_K\}$

```
for  $k = 1, \dots, K$  let  $\mathbf{r}(k)$  be a randomly chosen point from  $D$ ;  
while changes in clusters  $C_k$  happen do  
    form clusters:  
    for  $k = 1, \dots, K$  do  
         $C_k = \{\mathbf{x} \in D \mid d(\mathbf{r}_k, \mathbf{x}) \leq d(\mathbf{r}_j, \mathbf{x}) \text{ for all } j = 1, \dots, K, j \neq k\}$ ;  
    end;  
    compute new cluster centers:  
    for  $k = 1, \dots, K$  do  
         $\mathbf{r}_k$  = the vector mean of the points in  $C_k$   
    end;  
end;
```


Complexity

- K number of cluster centers
- n sample size
- I number of iterations
- *Suggestions?*


K-Means: Initial Centroids




K-Means: Initial Partition




K-Means: Iteration 2



K-Means: Iteration 20




Another Example: K-Means At Work




$K=2$

Arbitrarily choose K
object as initial
cluster center


Assign
each
objects
to most
similar
center




↑ reassign



Update
the
cluster
means

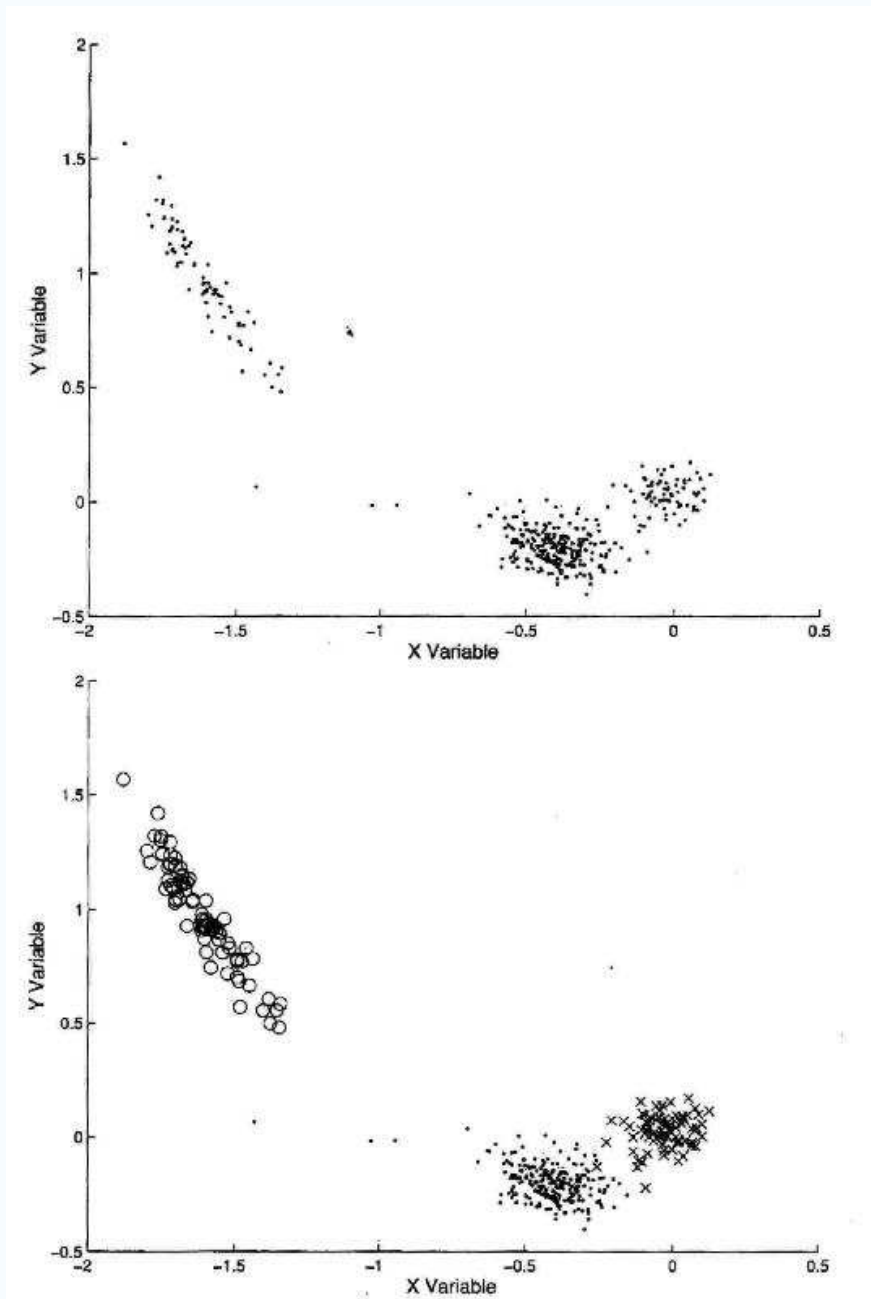


↓ reassign




Update
the
cluster
means

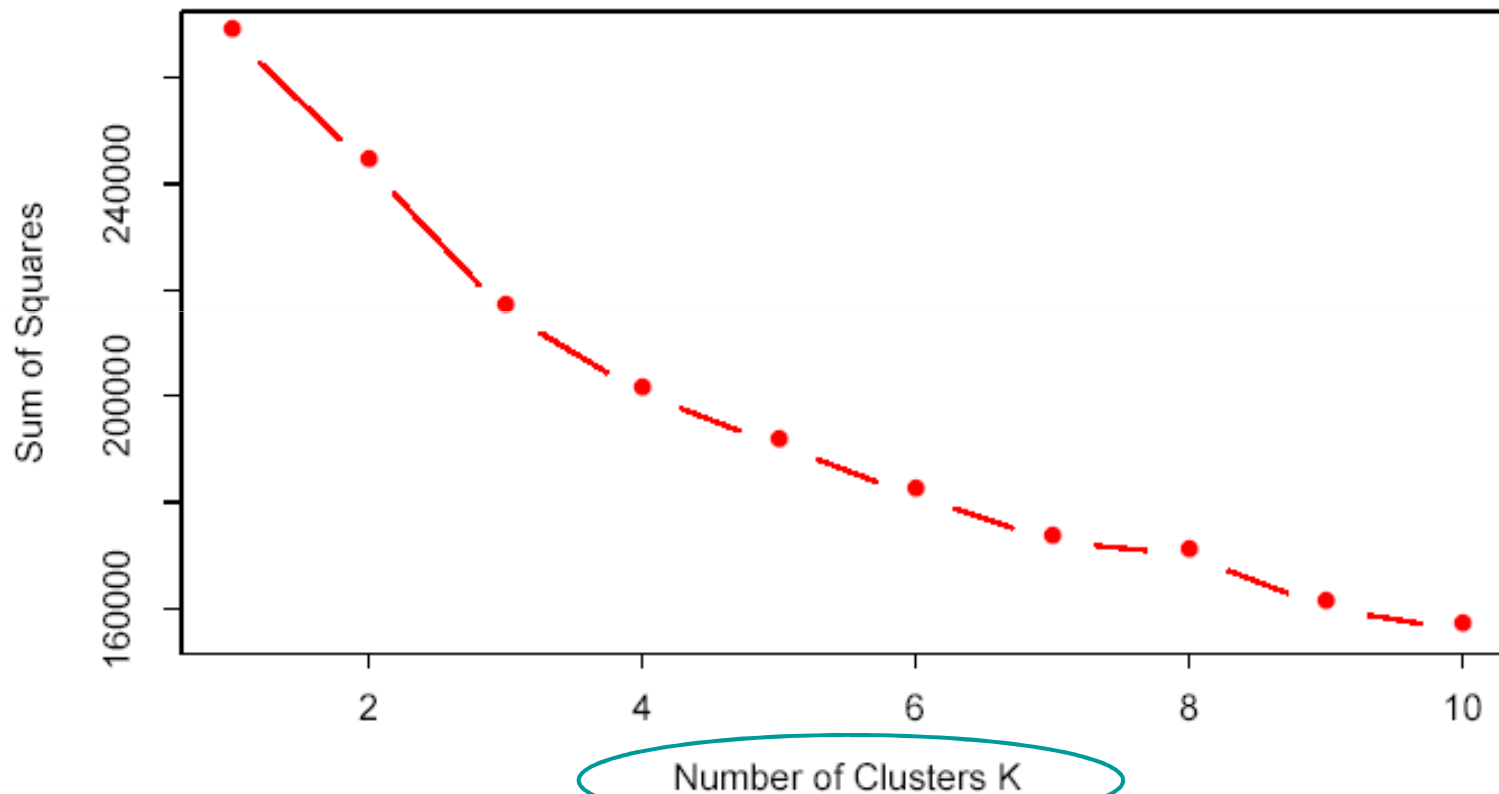
Antenna Data with/without Class Labels



Trajectory of Cluster Means on Antenna Data



K-Means Applied to Microarray Data



Y axis: total within-cluster sum of squares

K-Means Discussion 1

- Most suitable for real-valued data
- Prefers small, spherical clusters, unable to discover clusters with *non-convex shapes*
- Converges to a local optimum („trapped in local minimum“)
- *Not too many theoretical results on convergence*
- Example for local minimum:
 - four instances at the vertices of a two-dimensional rectangle
 - two cluster centers at the midpoints of the rectangle's long sides

K-Means Discussion 2

- Result can vary significantly based on initial choice of seeds
- Simple way to increase chance of finding a global optimum: restarts with different random seeds (“random restarts”)
- Linear in the examples - suitable for *large datasets*
- **K** has to be specified by the user - extensions as X-Means possible

Probability-Based/Model- Based Clustering: The EM Algorithm

Probability-Based Clustering


- Problem with previous approaches:
 - *ad hoc*: what are actually the clusters that are obtained? (“semantics”)
 - instances are “deterministically” assigned to one cluster
- From a probabilistic perspective, we want to find the *most likely clusters given the data*
- Also: instance only has certain *probability of belonging to a particular cluster*

Finite Mixtures

- Probabilistic clustering algorithms model the data as a mixture of distributions
- They are called *finite mixtures* because there is only a finite number of clusters being represented
- *Each cluster* is represented by *one distribution*
 - distribution governing the probabilities of attribute values in the clusters
- Usually individual distributions are normal distributions
- Distributions are combined using cluster weights

A Two-Class Mixture Model

| data | | | | | |
|-------|----|---|----|---|----|
| A | 51 | B | 62 | B | 64 |
| A | 43 | A | 47 | A | 51 |
| B | 62 | A | 52 | A | 52 |
| B | 64 | B | 64 | B | 62 |
| A | 45 | A | 51 | A | 49 |
| A | 42 | B | 65 | A | 48 |
| A | 46 | A | 48 | B | 62 |
| A | 45 | A | 49 | A | 43 |
| A | 45 | A | 46 | A | 40 |
| model | | | | | |
| A | 48 | A | 39 | A | 51 |
| B | 64 | B | 62 | A | 48 |
| A | 51 | B | 64 | B | 64 |
| B | 63 | A | 52 | A | 42 |
| A | 43 | B | 63 | A | 48 |
| B | 65 | B | 64 | A | 41 |
| B | 66 | A | 48 | | |
| B | 65 | B | 64 | | |
| A | 46 | A | 48 | | |



Using the Mixture Model

- The probability of an instance x belonging to cluster A is:

$$\Pr[A | x] = \frac{\Pr[x | A] \Pr[A]}{\Pr[x]} = \frac{f(x; \mu_A, \sigma_A) p_A}{\Pr[x]}$$

with

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- The *likelihood* of a single instance x given *all clusters* is:

$$\Pr[x | \text{the distributions}] = \sum_i \Pr[x | \text{cluster}_i] \Pr[\text{cluster}_i]$$

Learning the Clusters

- Assume we know that there are k clusters
- To learn the clusters we need to determine their *parameters*, i.e., their means and standard deviations
- We actually have a *performance criterion*: *likelihood of training data* given the clusters
- Fortunately, there exists an algorithm that finds a local maximum of the likelihood

EM Algorithm for Clustering

- EM algorithm: expectation-maximization algorithm
- In a sense, generalization of k-means to *probabilistic setting*
- Similar iterative procedure:
 1. Calculate cluster probability for each instance (Expectation step)
 2. Estimate distribution parameters based on the cluster probabilities (Maximization step)
- Cluster probabilities are stored as instance weights

More on EM

- Estimating parameters from weighted instances:

$$\mu_A = \frac{w_1 x_1 + w_2 x_2 + \dots + w_n x_n}{w_1 + w_2 + \dots + w_n}$$

$$\sigma_A^2 = \frac{w_1 (x_1 - \mu)^2 + w_2 (x_2 - \mu)^2 + \dots + w_n (x_n - \mu)^2}{w_1 + w_2 + \dots + w_n}$$

Log-Likelihood


- Procedure stops when *log-likelihood* saturates

$$\log \prod_{x_i \in D} \text{likelihood}(x_i) = \log \prod_{x_i \in D} \sum_{k \in \text{Cluster}} p_k P(x_i | C_k) =$$
$$\sum_{x_i \in D} \log \sum_{k \in \text{Cluster}} p_k P(x_i | C_k)$$

- For two clusters A and B:

$$\sum_{x_i \in D} \log(p_A P(x_i | A) + p_B P(x_i | B))$$

EM: Log-Likelihood over a Few Iterations



Special Case of EM: Putting it All Together

while log likelihood changes **do**

 % calculate cluster probability

 % for each instance (E-step)

for each instance x_i **do**

$$w_i = f(x_i; \mu_A, \sigma_A) p_A / P[x_i]$$

 % estimate distribution parameters based

 % on the cluster probabilities (M-step)

$$\mu_A = \sum w_i x_i / \sum w_i$$

$$\sigma_A^2 = \sum w_i (x_i - \mu_A)^2 / \sum w_i$$

 % analogously for μ_B and σ_B^2 with $(1-w_i)$

Special Case of EM: Putting it All Together

while log likelihood change
% calculate cluster proba
% for each instance (E-step)
for each instance x_i **do**

$$w_i = f(x_i; \mu_A, \sigma_A) p_A / P[x_i]$$

% estimate distribution parameters based
% on the cluster probabilities (M-step)


$$\mu_A = \sum w_i x_i / \sum w_i$$

$$\sigma_A^2 = \sum w_i (x_i - \mu_A)^2 / \sum w_i$$

% analogously for μ_B and σ_B^2 with $(1-w_i)$

```
for k = 1, ..., K let  $r(k)$  be a randomly chosen point from  $D$ ;  
while changes in clusters  $C_k$  happen do  
  form clusters:  
  for k = 1, ..., K do  
     $C_k = \{x \in D \mid d(r_k, x) \leq d(r_j, x) \text{ for all } j = 1, \dots, K, j \neq k\}$ ;  
  end;  
  compute new cluster centers:  
  for k = 1, ..., K do  
     $r_k$  = the vector mean of the points in  $C_k$   
  end;  
end;
```



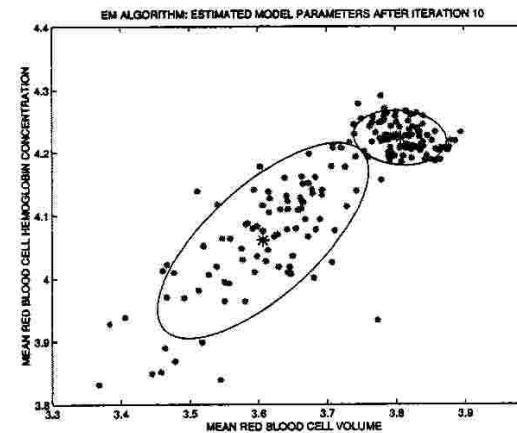

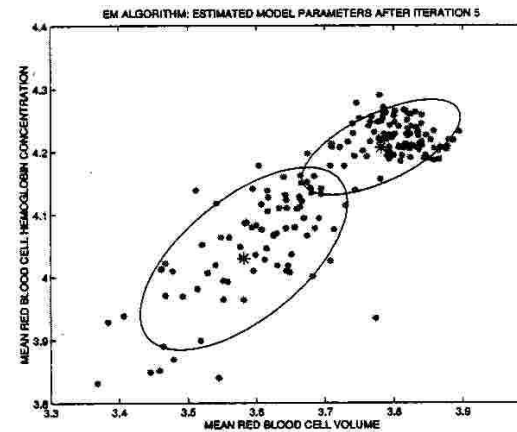
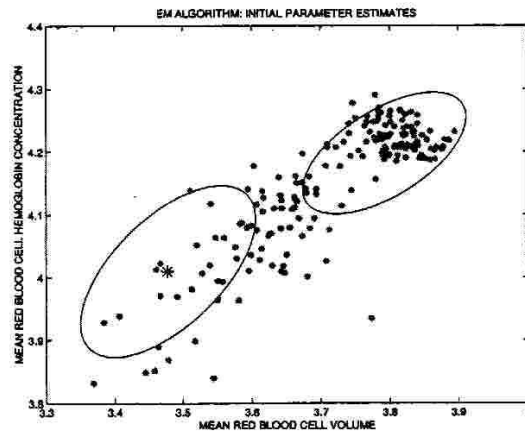

Red Blood Cell Data



Circles:
healthy,
crosses:
iron
deficient
anemia

Example Run EM on Red Blood Cell Data

3σ covariance
ellipses and
means for
iterations
1, 2, 5, 10, 25



Extensions 1

- Using more than two distributions: easy
- Several attributes: easy if independence is assumed
- Correlated attributes: difficult
 - Modeled jointly using a bivariate normal distribution with a covariance matrix
 - With n attributes this requires estimating $n + n(n+1)/2$ parameters
- Nominal attributes: easy if independent

Extensions 2

- Correlated nominal attributes: difficult
 - two correlated attributes result in $v_1 \times v_2$ parameters
- Missing values: easy
- Distributions other than the normal distribution can be used:
 - log-normal if predetermined minimum is given
 - log-odds if bounded from above and below
 - Poisson for attributes that are integer counts
- Cross-validation can be used to estimate k

Discussion

- Can be used to fill in missing values
- Big advantage of probabilistic clustering schemes: likelihood of data can be estimated and used to compare different clusterings