



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

# Data Mining – Relational Learning

Andreas Karwath    Jörg Wicker

Johannes Gutenberg University Mainz

January 26, 2017

# Acknowledgments

- Slides partially taken from
  - Stefan Kramer
  - Tom Mitchel
  - Johannes Frnkranz

# Section 1


## Introduction

# The Story so far ...

## Item mining



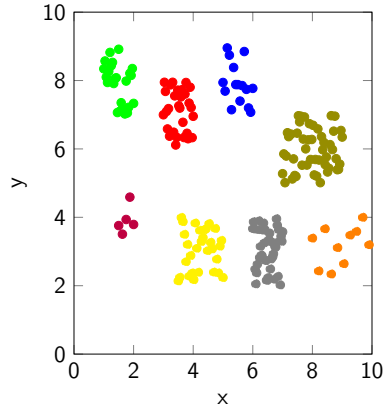
## Graph mining

- Subgraph pattern  $p_2$ : 
- Database  $D$  consisting of graphs  $b)$ ,  $c)$ , and  $d)$ :



- $\text{support}(p_2, D) = 2$

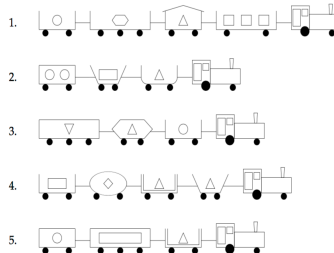
## Clustering



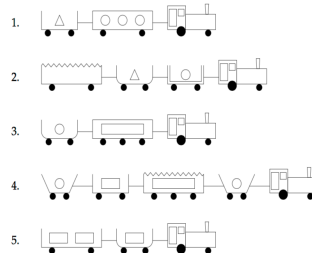
# Relational Learning Introduction I

- Offers a versatile way of describing data and patterns (pattern language)
- Can be applied for item set mining as well as for graph mining or ...
- just for all sorts of relational problems.
- But, what are relational problems?
- For example the train example from last weeks outlook:

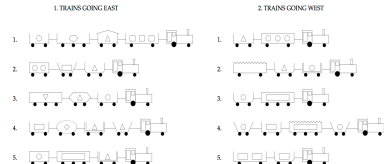
1. TRAINS GOING EAST



2. TRAINS GOING WEST



# Relational Learning Introduction II



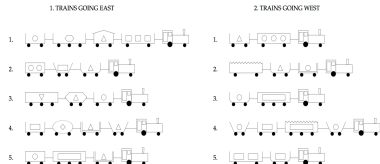
- So what rule can describe the difference between east and west trains?
- One possible explanation (or so-called hypothesis):  
Trains go east if they possess a short car with a roof.
- The rule is fine for humans, but how can one model this using computers?

- One way would be to “flatten” it:

#Ex.	$C1_{shape}$	$C1_{axes}$	$C1_{roof}$	$C1_{\#objects}$	$C1_{objectShape}$	...	Class
1	rectangle	2	open	3	square	...	east
2	bucket	2	open	1	triangle	...	east
...	...	...	...	...	...	...	...

- However, this is not ideal:
  - change in the order of cars
  - different shape in the cargo
  - varying number of cars
  - ...

# Relational Learning Introduction III



- A more “natural” way of describing these examples is in language able to cater for relational data: e.g. Prolog

- Example:

```
east(e1).  
has_car(e1,e1c1).  
shape(e1c1,rectangle).  
length(e1c1,short).  
roof(e1c1,open).  
carry(e1c1,e1c1o1,square).  
carry(e1c1,e1c1o2,square).  
carry(e1c1,e1c1o3,square).  
...
```

## Quick Prolog Recap

- Variables: `X` , `Y` , `A` , `B` , `Train` , `Car`
- Terms: `square` , `t1` , `1` , `[1,2,3]`
- Predicates: `east/1` , `shape/1` , `carry/3`

- Facts:

```
shape(e1c1,rectangle).  
carry(e1c1,e1c1o1,square).
```

- Rules:

```
east(Train) :- has_car(Train,Car), length(Car,short), roof(Car,open).
```



# Logical reasoning: Deduction

- From rules to facts

$$B \cup T \vdash E$$

- $B$ : Background knowledge
- $T$ : Theory
- $E$ : Examples

$B$	$\cup$	$T$	$\vdash$	$E$
-----	--------	-----	----------	-----

```
mother(penelope,victoria).
mother(penelope,arthur).
father(christopher,victoria).
father(christopher,arthur).
```

```
parent(X,Y) :- father(X,Y).
parent(X,Y) :- mother(X,Y).
```

```
parent(penelope,victoria).
parent(penelope,arthur).
parent(christopher,victoria).
parent(christopher,arthur).
```

# Logical reasoning: Induction

- From facts to rules

$$B \cup E \vdash T$$

- $B$ : Background knowledge
- $T$ : Theory
- $E$ : Examples

$B$	$\cup$	$E$	$\vdash$	$T$
-----	--------	-----	----------	-----

```
mother(penelope,victoria).
mother(penelope,arthur).
father(christopher,victoria).
father(christopher,arthur).
```

```
parent(penelope,victoria).
parent(penelope,arthur).
parent(christopher,victoria).
parent(christopher,arthur).
```

```
parent(X,Y) :- father(X,Y).
parent(X,Y) :- mother(X,Y).
```

# Induction of a classifier

- Given:
  - background knowledge  $B$
  - a set of training examples  $E$
  - a classification  $c \in C$  for each example  $e$
- Find: a theory  $T$  (or *hypothesis*) such that  $B \cup T \vdash c(e)$ , for all  $e \in E$

## Induction of a classifier: Train Example

- $B$ : relations `has_car` and *car properties* (length, roof, shape, etc.)  
example.: `has_car(t1,c11), shape(c11,bucket)`
- $E$ : the trains `t1` to `t10`
- $C$ : east, west (or  $\neg$ east)
- Possible  $T$ :  
`east(Train) :- has_car(Train,Car), length(Car,short), roof(Car,open).`

# Learning as search

- Given:
  - Background knowledge  $B$
  - Theory Description Language  $T$  (logic)
  - Positives examples  $P$  (class  $\oplus$ )
  - Negative examples  $N$  (class  $\ominus$ )
  - A covering relation  $\text{covers}(B, T, e)$
- Find: a theory that covers
  - all positive examples (completeness)
  - no negative examples (consistency)

# Learning as search

- Covering relation:

$$\text{covers}(B, T, e) \Leftrightarrow B \cup T \vdash e$$

- A theory is a set of rules

- a rule  $R_i$  is in *CNF* ( in the form:  $l_1 \wedge l_2 \wedge \dots l_n$  or  $\{l_1, l_2, \dots, l_n\}$ )
- the complete theory  $T$  in *DNF* (in the form  $R_1 \vee R_2 \vee \dots R_m$  or  $\{R_1; R_2, \dots; R_m\}$ )

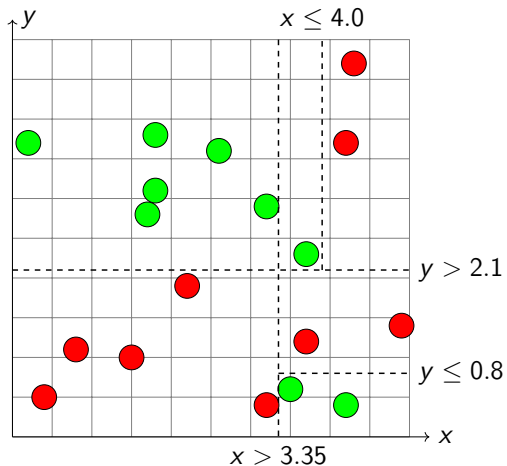
- Each rule is searched separately (efficiency)
- A rule must be consistent (cover no negatives), but not necessary complete
- Separate-and-conquer strategy
- Remove from  $P$  the examples already covered

# Separate-and-Conquer Strategy - Rule Learning Example

- Example (using propositional logic with just 2 numeric dimensions)

Ex. #	X	Y	Class
1	0.4	0.5	$\ominus$
2	4.3	4.7	$\ominus$
3	2.2	1.9	$\ominus$
4	0.8	1.1	$\ominus$
5	1.5	1.0	$\ominus$
6	3.2	0.4	$\ominus$
7	4.2	3.7	$\ominus$
8	3.7	1.2	$\ominus$
9	4.9	1.4	$\ominus$
10	1.8	3.8	$\oplus$
11	2.6	3.6	$\oplus$
12	3.2	2.9	$\oplus$
13	1.8	3.1	$\oplus$
14	3.5	0.6	$\oplus$
15	4.2	0.4	$\oplus$
16	0.2	3.7	$\oplus$
17	1.7	2.8	$\oplus$
18	3.7	2.3	$\oplus$

# Separate-and-Conquer Strategy - Rule Learning Example





## Separate-and-Conquer System - pFOIL

- Propositional variant of FOIL (Quinlan, 1993), one of the earliest and simplest relational learning system (*inductive logic programming* (ILP) - machine learning in predicate logic)
- Search heuristic: weighted information gain
- Search strategy: hill climbing
- Stopping criterion: encoding length restriction

# pFOIL - Algorithm

---

```
1  $Pos \leftarrow$  positive examples;
2  $Neg \leftarrow$  negative examples;
3  $LearnedRules \leftarrow \{\}$ ;
4 while  $Pos \neq \{\}$  do
5    $NewRule \leftarrow$  most general rule ( $\top$ );
6    $NewRuleNeg \leftarrow Neg$ ;
7   while  $NewRuleNeg \neq \{\}$  do
8      $NewRule \leftarrow \operatorname{argmax}_{NewRule' \in \rho_s(NewRule)} pFOILGain(NewRule')$ ;
9      $NewRuleNeg \leftarrow \operatorname{cover}(B, NewRule, NewRuleNeg)$ ;
10    //  $NewRuleNeg$  = subset of  $NewRuleNeg$  satisfying  $NewRule$ 
11  end
12   $LearnedRules \leftarrow \operatorname{append}(LearnedRules, (NewRule))$ ;
13   $Pos \leftarrow Pos \setminus \operatorname{cover}(B, NewRule, Pos)$ ;
14 end
15 return  $LearnedRules$ 
```

---

## pFOIL - Information Gain (pFOILGain)

$$pFOILGain(R \wedge L) = p_1 \left( \log_2 \frac{p_1}{p_1 + n_1} - \log_2 \frac{p_0}{p_0 + n_0} \right)$$

■ where

- $L$  = candidate literal to be added to rule  $R$
- $p_0$  = number of positive examples of  $R$
- $n_0$  = number of negative examples of  $R$
- $p_1$  = number of positive examples of  $R \wedge L$
- $n_1$  = number of negative examples of  $R \wedge L$

## pFOIL - Stopping Criterion Based on Encoding Length Restriction

- Minimum Description Length (MDL) principle
- Training set of size  $|T|$ , a rule accounts for  $p$  positive examples
- Number bits required to identify (choose) those examples
- Versus coding length per literal:  $1 + \log_2(|Literals|)$
- Tries to avoid learning complicated rules (covering only a few examples) ensuring that number of bits needed to encode a rule (clause)  $<$  number of bits needed to encode the instances covered by it.

# Refinement Operator $\rho$ - Propositional Case

Rules are in conjunctive normal form (CNF):

- In the propositional case (examples are given in form of *attributes* and *values*), literals are just in the form *attribute = value*.
- Example of  $\rho$ :
  - $R_0 = \top$
  - $R_1 \in \rho(R_0) = \{l_1, l_2, \dots, l_n\}$ , e.g.  $R_1 = l_1$
  - $R_2 \in \rho(R_1) = \{l_1 \wedge l_2, l_1 \wedge l_3, \dots, l_1 \wedge l_n\}$
  - ...
- i.e.  $\rho(R)$  specializes  $R$  by adding a literal to the conjunction

# Refinement Operator $\rho$ - First-Order (FOIL) Case

Rules are in clausal normal form (CNF):

- The most general rule is the Horn clause using the predicate we want to learn, i.e.  
 $\text{east}(T) :- .$
- $\rho(R)$  specializes a rule  $R$  by either:
  - adding  $p(V_1, \dots, V_n)$  where  $p$  is a predicate and  $V_i$  are variables not yet occurring in the rule(clause)  $R$ .
  - adding  $\text{equal}(V_j, V_k)$  or its **negation**, where  $V_j$  and  $V_k$  are variables already present in the rule
- example:  
 $\rho(\text{east}(T) :- .) = \{\text{east}(T) :- \text{has\_car}(A, B) ., \text{east}(T) :- \text{roof}(A, B) ., \dots\}$
- or  $\rho(\text{east}(T) :- \text{has\_car}(A, B) .) =$   
 $\{\text{east}(T) :- \text{has\_car}(A, B), \text{roof}(C, D) ., \text{east}(T) :- \text{has\_car}(A, B), \text{equal}(T, A) ., \dots\}$

# Refinement Operator $\rho$ - General Overview

- Item set mining:
  - $\rho(I)$ : add an item to current item set  $I$
- Graph mining (data set of graph setting):
  - $\rho(G)$ : add an edge to current sub graph  $G$
- Learning rules (propositional):
  - $\rho(R_p)$ : add a literal to current rule  $R_p$
- Learning rules (first-order logic):
  - $\rho(R_f)$ :
    - add a literal to current rule  $R_f$  or
    - apply an elementary substitution to  $R_f$ .

# General Pattern Mining in Relation Data

- The idea of this refinement operator (and a bit more sophisticated ones) can also be applied to pattern mining.
- WARMR is one example of APRIORI in predicate logic, others are CARMR, *etc.*
- However, the search space explodes if one is not careful:
  - use of so-called language bias
  - use of condensed representations
  - ...



# Outlook

- ProbLog
- relational learning and probabilities